

Parallel K-means Algorithm Using MPI

Fan Xie fx349@nyu.edu

1. Clustering Problem Description

1.1 Description

Given a set of data points $X = \{x_1, x_2, \dots, x_n\}$, where each data point is a d -dimensional real vector, clustering algorithm aims to partition those data points into k sets $C = \{C_1, C_2, \dots, C_k\}$ so as to minimize the within-cluster sum of squares. Each point is closer to points inside than outside of its subset[1].

1.2 Mathematical Definition:

$$\begin{aligned} J_{k\text{-means}}(C_1, \dots, C_k) &= \sum_{i=1}^k \sum_{x \in C_i} d(x, \mu(C_i))^2 \\ &= \min_{\mu_1, \dots, \mu_k \in \mathcal{X}} \sum_{i=1}^k \sum_{x \in C_i} d(x, \mu_i)^2 \end{aligned}$$

2. K-means Algorithm Description

2.1 Description

Finding the best Partitions in a clustering problem is actually an NP-hard problem[1], so we need some kind of approximation algorithm to solve it in

practice. K-means is one of the heuristic algorithms to solve the clustering problem. It converges quickly to a local optimum and its performance heavily depends on the initialization step. So usually we run K-means algorithm several times to get a better result.

2.2 Algorithm

Input: *a set of data points $X = \{x_1, x_2, \dots, x_n\}$ and a parameter k*

Initialize: *Randomly choose initial k centroids*

Repeat until convergence:

Assignment step: *Compute the distance from each point x_i to each cluster centroid C_j and assign each point to the centroid it is closest to*

Update step: *Recompute each centroid as the mean of all points assigned to it*

3. Implement K-means Algorithm Using MPI

3.1 Description

As the number of data points grows, the k-means algorithm can be very slow. So we need to find some way to parallel it. As we can see, the assignment step is the most time-cost step in this algorithm. In the assignment step, each data point does their own things individually. and no information except the k centers needs to be shared[2]. So we can easily adopt a data parallelism method to speed up the process.

3.2 Algorithm

Initialize:

1. Master process read all the data points and randomly initialize k centroids.
2. Master process divide those data points evenly and send them to the corresponding processes (MPI_Bcast, MPI_Scatter)

Repeated until converge:

Assignment step

1. Master process broadcast k centroids to all the processes (MPI_Bcast)
2. Each process calculate new centroids for all the data points that belong to it

Update Step

3. Each process tells the master process what's the new centroid of each data point (MPI_Gather)
4. Master process recompute k new centroids base on the new assignment

4. Code Repository

<https://github.com/xiefan46/homework2/tree/master/project/k-means>

5. Reference

[1] https://en.wikipedia.org/wiki/K-means_clustering

[2] http://www.goldsborough.me/c++/python/cuda/2017/09/10/20-32-46-exploring_k-means_in_python,_c++_and_cuda/